

## WHAT IS PYTORCH?

### Definition

PyTorch is an open-source deep learning framework by Meta AI providing:

- (1) N-dimensional tensor computation (GPU NumPy)
- (2) Automatic differentiation (autograd)
- (3) Dynamic computation graph (Define-by-Run)

### Why Dynamic Graph?

Graph built *on the fly* during execution. Use normal Python `if/else`, `print()`, `pdb` directly on models. TF1.x used static graphs — harder to debug.


## TENSOR — THE DNA

### Definition

A **Tensor** is a generalized multi-dimensional array unifying scalars, vectors, and matrices under one structure.

### Tensor Dimensionality

Name	Dim	Notation	Example
Scalar	0D	$x \in \mathbb{R}$	5
Vector	1D	$\mathbf{x} \in \mathbb{R}^n$	[1, 2, 3]
Matrix	2D	$\mathbf{X} \in \mathbb{R}^{m \times n}$	grid
3D Tensor	3D	$\mathcal{T} \in \mathbb{R}^{d \times m \times n}$	RGB image
4D Tensor	4D	$\mathcal{T} \in \mathbb{R}^{B \times C \times H \times W}$	batch



### Batch Image Shape

$\mathcal{T} \in \mathbb{R}^{B \times C \times H \times W}$ ;  $B=32$  (batch),  $C=3$  (RGB),  $H, W=224$

## CREATING TENSORS

Syntax	Purpose
<code>torch.tensor([1,2,3])</code>	From Python list
<code>torch.zeros(3,4)</code>	All zeros
<code>torch.ones(3,4)</code>	All ones
<code>torch.rand(3,4)</code>	Uniform $\mathcal{U}[0, 1]$
<code>torch.randn(3,4)</code>	Normal $\mathcal{N}(0, 1)$
<code>torch.arange(0,10,2)</code>	Integer range
<code>torch.linspace(0,1,5)</code>	Evenly spaced
<code>torch.eye(3)</code>	Identity $\mathbf{I}_3$
<code>torch.zeros_like(t)</code>	Same shape, zeros
<code>torch.empty(2,3)</code>	Uninitialized

### Key Distinction

`torch.tensor()` copies data.  
`torch.as_tensor()` shares memory with NumPy arrays.

## TENSOR ATTRIBUTES (BIG THREE)

### Three Core Properties

<code>t.dtype</code>	data type stored
<code>t.shape</code>	dimensions ( <code>torch.Size</code> )
<code>t.device</code>	cpu / cuda / mps
Also: <code>t.ndim</code> , <code>t.numel()</code>	

dtype	Use
<code>float32</code>	Default NN weights
<code>float16</code>	GPU half-precision
<code>int64</code>	Default integers
<code>bool</code>	Masks, conditions

## DEVICE MANAGEMENT

```
device = "cuda" if torch.cuda.is_available()
() else "cpu"
```

```
t = t.to(device)      # move tensor
t = t.cuda()          # shorthand GPU
t = t.cpu()          # back to CPU
```

### Common Error

Two tensors must be on **same device**. Mixing cpu/cuda throws `RuntimeError`. Move **both** model **and** data!

## TENSOR OPERATIONS

### Element-wise

$(a \odot b)_{ij} = a_{ij} \cdot b_{ij}$  — operates at each position independently

```
a*b; a+b; a-b; a/b
torch.sqrt(a); torch.exp(a); torch.log(a)
```

### Matrix Multiplication — MOST IMPORTANT

For  $\mathbf{A} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{B} \in \mathbb{R}^{n \times p}$ :

$$\mathbf{C} = \mathbf{AB} \in \mathbb{R}^{m \times p}, \quad C_{ij} = \sum_{k=1}^n A_{ik}B_{kj}$$

```
A @ B           # preferred (batched
too)
torch.matmul(A,B) # same
torch.mm(A,B)   # 2D only
```

$\mathbf{a} * \mathbf{b} = \text{element-wise}$        $\mathbf{a} @ \mathbf{b} = \text{matmul}$

### Aggregation

```
t.sum(); t.mean(); t.std()
t.max(); t.argmax(); t.argmin()
t.sum(dim=0) # collapse rows
t.sum(dim=1) # collapse columns
```

$$\mu = \frac{1}{n} \sum_i x_i \quad \sigma = \sqrt{\frac{1}{n} \sum_i (x_i - \mu)^2}$$

## SHAPE MANIPULATION


Operation	Effect
<code>reshape(r,c)</code>	New shape, may copy
<code>view(r,c)</code>	New shape, shares mem
<code>squeeze()</code>	Remove size-1 dims
<code>unsqueeze(0)</code>	Add dim at pos 0
<code>permute(2,0,1)</code>	Reorder all dims
<code>transpose(0,1)</code>	Swap two dims
<code>flatten()</code>	Collapse to 1D
<code>cat([a,b],dim=0)</code>	Concat existing dim
<code>stack([a,b],dim=0)</code>	Create new dim

### The -1 Trick

`t.reshape(32,-1)`: PyTorch infers the second dim. If  $t$  has  $32 \times N$  elements, result is  $(32, N)$ .

### PIL vs PyTorch Format

PIL/NumPy:  $(H, W, C)$   
 PyTorch:  $(C, H, W)$   
 Convert: `tensor.permute(2,0,1)`



## AUTOGRAD — AUTO DIFFERENTIATION


### Definition

**Autograd** records operations on tensors with `requires_grad=True`, builds a DAG, and computes gradients via backpropagation when `.backward()` is called.

$$\frac{dy}{dx} = 9x^2 + 2 \quad \text{At } x=4: 9(16) + 2 = 146$$

```
x = torch.tensor(4.0, requires_grad=True)
y = 3*x**3 + 2*x + 1
y.backward()
print(x.grad) # tensor(146.)
```

### Computation Graph



```
x = torch.tensor(3.0, requires_grad=True)
y = x ** 2
y.backward()
print(x.grad) # 6.0 (= 2*3)
```

```
# Stop tracking (inference)
with torch.no_grad():
    pred = model(x)
    t.detach() # new tensor, no grad history
```

### Sacred Training Loop Order

1. `pred=model(X)` forward
2. `loss=criterion(pred,y)` loss
3. `optimizer.zero_grad()` CLEAR
4. `loss.backward()` backprop
5. `optimizer.step()` update

### Chain Rule — Foundation of Backprop

If  $\mathcal{L} = f(g(h(\mathbf{x})))$ :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \underbrace{\frac{\partial \mathcal{L}}{\partial f}}_{\text{loss}} \cdot \underbrace{\frac{\partial f}{\partial g}}_{\text{L3}} \cdot \underbrace{\frac{\partial g}{\partial h}}_{\text{L2}} \cdot \underbrace{\frac{\partial h}{\partial \mathbf{x}}}_{\text{L1}}$$

### Worked Example

$y = 3x^3 + 2x + 1$ , find  $dy/dx$  at  $x = 4$ :

### Why zero\_grad Matters

Gradients **accumulate** (add to `.grad`). Without clearing, batch 2's gradients pile on batch 1's — updates are wrong from iteration 2 onward.

**nn.Module — THE BACKBONE****Definition**

`nn.Module` is the base class for all models. Every custom model **must** inherit from it. Implement: `__init__` (layers) and `forward` (data flow).

```
import torch.nn as nn

class MyNet(nn.Module):
    def __init__(self):
        super().__init__() # ALWAYS first!
        self.fc1 = nn.Linear(784, 128)
        self.fc2 = nn.Linear(128, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.fc1(x))
        return self.fc2(x)

model = MyNet()
```

**KEY LAYERS — MATH****nn.Linear — Fully Connected**

$$\mathbf{y} = \mathbf{x}\mathbf{W}^\top + \mathbf{b}$$

$$\mathbf{W} \in \mathbb{R}^{\text{out} \times \text{in}}, \mathbf{b} \in \mathbb{R}^{\text{out}}$$

Params = ( $\text{in} \times \text{out}$ ) +  $\text{out}$


$$\text{Linear}(512, 256): 512 \times 256 + 256 = 131,328$$

**Why Activations Are Needed****Key Insight**

Without non-linearity:  $\mathbf{W}_2(\mathbf{W}_1\mathbf{x}) = (\mathbf{W}_2\mathbf{W}_1)\mathbf{x}$  — stacking linear layers  $\equiv$  one linear layer. Activations enable complex patterns.


**ReLU**

$$f(x) = \max(0, x), \quad f'(x) = \begin{cases} 1 & x > 0 \\ 0 & x \leq 0 \end{cases}$$

**Sigmoid and Tanh**

$$\sigma(x) = \frac{1}{1 + e^{-x}} \in (0, 1), \quad \sigma' = \sigma(1 - \sigma)$$

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \in (-1, 1)$$

**Softmax**

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}, \quad \sum_i p_i = 1$$

**Activation Summary**

Name	Range	Use	Vanish?
ReLU	$[0, \infty)$	Hidden	No
Sigmoid	$(0, 1)$	Binary out	Yes
Tanh	$(-1, 1)$	Hidden, RNN	Mild
Softmax	$(0, 1)^K$	Multi-class	—

**LOSS FUNCTIONS****MSE — Regression**

$$\mathcal{L}_{\text{MSE}} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

`nn.MSELoss()` — linear output.

**Binary Cross Entropy**

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{n} \sum_i [y_i \log \hat{y}_i + (1-y_i) \log(1-\hat{y}_i)]$$

**Use BCEWithLogitsLoss**

Applies sigmoid internally. More numerically stable. Never use `BCELoss` + manual sigmoid.

**Cross Entropy — Multiclass**

$$\mathcal{L}_{\text{CE}} = -\frac{1}{n} \sum_i \sum_j y_{ij} \log \hat{y}_{ij}$$


**CRITICAL — No Softmax Before CE**

`CrossEntropyLoss`=`LogSoftmax`+`NLLLoss`. Passing softmax outputs applies softmax twice — **silent wrong results**.

**Loss Selection — Memorise**

Task	Loss	Output
Regression	<code>MSELoss</code>	linear
Binary cls	<code>BCEWithLogitsLoss</code>	linear
Multi-class	<code>CrossEntropyLoss</code>	linear

## BACKPROPAGATION — VISUAL



**Forward:** data flows input → output, graph built.

**Backward:** gradients flow output → input via chain rule.

$$\text{Each weight: } w \leftarrow w - \eta \frac{\partial \mathcal{L}}{\partial w}$$

## OPTIMIZERS

### ■ SGD

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta g_t$$

With momentum:  $v_t = \beta v_{t-1} + (1 - \beta)g_t$

```
torch.optim.SGD(model.parameters(), lr
                =0.01)
torch.optim.SGD(model.parameters(), lr
                =0.01, momentum=0.9)
```

### ■ Adam — Adaptive Moment Estimation

## Adam Equations

$$\begin{aligned} m_t &= \beta_1 m_{t-1} + (1 - \beta_1)g_t && \text{(1st moment)} \\ v_t &= \beta_2 v_{t-1} + (1 - \beta_2)g_t^2 && \text{(2nd moment)} \\ \hat{m}_t &= \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} && \text{(bias corr)} \\ \mathbf{w} &\leftarrow \mathbf{w} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}} \end{aligned}$$

Defaults:  $\beta_1=0.9$ ,  $\beta_2=0.999$ ,  $\epsilon=10^{-8}$

```
torch.optim.Adam(model.parameters(), lr=1e-3)
```

Optimizer	Best For	lr
SGD	Vision tasks	0.01
SGD+momentum	Faster conv.	0.01
Adam	Default choice	0.001
AdamW	Transformers	0.001

## BATCH NORMALISATION

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}, \quad y_i = \gamma \hat{x}_i + \beta$$

$\gamma, \beta$  are learned. During eval(): uses running stats.

```
nn.BatchNorm1d(num_features) # after Linear
nn.BatchNorm2d(num_channels) # after Conv2d
```

**Benefits:** faster training, higher LR toler-


ance, mild regularisation.

## DROPOUT

$$\tilde{h}_i = \frac{h_i \cdot m_i}{1 - p}, \quad m_i \sim \text{Bernoulli}(1 - p)$$

Randomly zeros neurons with prob  $p$  during training.

```
nn.Dropout(p=0.5) # 50% zeroed (train)
# eval(): automatically disabled
```



## nn.Sequential

```
model = nn.Sequential(
    nn.Linear(784, 256),
    nn.BatchNorm1d(256),
    nn.ReLU(),
    nn.Dropout(0.3),
    nn.Linear(256, 128),
    nn.ReLU(),
    nn.Linear(128, 10)
)
```

## Sequential vs Module

**Sequential:** simple linear stacks.

**Module:** skip connections, multiple inputs/outputs, custom logic.

## DATASETS & DATALOADERS

### Definition

**Dataset:** Abstract class. Must implement `_len_` and `_getitem_`.  
**DataLoader:** Wraps Dataset; batching, shuffling, parallel loading.

```
from torch.utils.data import Dataset, DataLoader

class MyDataset(Dataset):
    def __init__(self, X, y):
        self.X, self.y = X, y

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]

ds = MyDataset(X_tensor, y_tensor)
dl = DataLoader(ds, batch_size=32, shuffle=True)
```

Param	Meaning
batch_size	Samples per iteration
shuffle=True	Randomise each epoch
num_workers	Parallel loading threads
drop_last	Drop final small batch

## MODEL MODES

### Always Switch Modes

`model.train()`: Dropout active, BN uses batch stats.  
`model.eval()`: Dropout off, BN uses running stats.  
 Forgetting this **silently** degrades performance!

```
model.train()
for X, y in train_dl:
    ... # training loop

model.eval()
with torch.no_grad():
    for X, y in val_dl:
        pred = model(X)
```

## CONVOLUTIONAL LAYERS (CNN)

### Output Spatial Size

$$H_{out} = \left\lceil \frac{H_{in} + 2P - K}{S} \right\rceil + 1$$


$P$ =padding,  $K$ =kernel,  $S$ =stride. Same for  $W_{out}$ .

### Parameter Count

$$\text{params} = C_{out} \times (C_{in} \times K^2 + 1)$$

`Conv2d(3, 64, 3): 64 × (3 × 9 + 1) = 1,792`

```
nn.Conv2d(3, 64, kernel_size=3, padding=1)
# padding=1 with kernel=3 -> same spatial
# size
nn.MaxPool2d(kernel_size=2, stride=2)
# halves both H and W
```



## SAVING & LOADING

```
torch.save(model.state_dict(), "model.pth")
)

model = MyNet()
model.load_state_dict(torch.load("model.pth"))
model.eval()
```

### Why state\_dict?

Saves only parameter tensors. Portable across codebases. Saving whole model pickles the class — breaks on refactor or rename.

## COMPLETE TRAINING PIPELINE

```
import torch, torch.nn as nn
from torch.utils.data import DataLoader, TensorDataset

# Data
X = torch.randn(1000, 20)
y = torch.randint(0, 3, (1000,))
ds = TensorDataset(X, y)
train_ds, val_ds = torch.utils.data.random_split(
    ds, [800, 200])
train_dl = DataLoader(train_ds, batch_size=32, shuffle=True)
val_dl = DataLoader(val_ds, batch_size=32)

# Model
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(20, 64), nn.ReLU(),
            nn.Linear(64, 32), nn.ReLU(),
            nn.Linear(32, 3))
    def forward(self, x): return self.net(x)

device = "cuda" if torch.cuda.is_available()
() else "cpu"
model = Net().to(device)
crit = nn.CrossEntropyLoss()
optim = torch.optim.Adam(model.parameters(),
(), lr=1e-3)

# Loop
for epoch in range(10):
    model.train(); tloss = 0
    for Xb, yb in train_dl:
        Xb, yb = Xb.to(device), yb.to(
            device)
        loss = crit(model(Xb), yb)
        optim.zero_grad(); loss.backward()
        optim.step(); tloss += loss.item()
    model.eval(); correct = 0
    with torch.no_grad():
        for Xb, yb in val_dl:
            Xb, yb = Xb.to(device), yb.to(
                device)
            correct += (model(Xb).argmax(
                1)==yb).sum().item()
    print(f"Ep{epoch+1} Loss:{tloss/len(
        train_dl):.3f} Acc:{correct /
    200:.3f}")

torch.save(model.state_dict(), "model.pth")
```

## PRACTICE QUESTIONS

### Section A — Conceptual

- Q1.** Explain dynamic vs static computation graphs. Why does PyTorch's approach aid debugging?
- Q2.** What are the **three** core tensor attributes? Describe each.
- Q3.** Why must `optimizer.zero_grad()` be called before `loss.backward()`?
- Q4.** Why should you *never* apply softmax before `nn.CrossEntropyLoss()`?
- Q5.** Difference between `model.train()` and `model.eval()`? Name 2 affected layers.
- Q6.** Write chain rule for  $\mathcal{L} = f(g(h(\mathbf{x})))$  w.r.t.  $\mathbf{x}$ .
- Q7.** What does `requires_grad=True` do? When is the graph built?

### Section B — Shape Tracing

- Q8.** `t=randn(4,3,28,28)`. Shapes after: `view(4,-1)`, `permute(0,2,3,1)`, `[:,0,:,:]`?
- Q9.** `a=randn(32,64)`, `b=randn(64,128)`. Shape of `a @ b`?
- Q10.** `t=randn(8,16)`, shape of `t.unsqueeze(0).unsqueeze(2)?` `view(4,-1)`:  $(4, 2352)$  since  $3 \times 28^2 = 2352$ ; `permute(0,2,3,1)`:  $(4, 28, 28, 3)$ ; `[:,0,:,:]`:  $(4, 28, 28)$
- Q11.** `t=randn(3,1,5)`, shape of `t.squeeze()`?
- Q12.** Shape  $(100, ) \rightarrow (100, 1)$ . What single call?

### Section C — Debug the Code

- Q13.** User wants matmul but writes `result=x*y` for  $(10, 5)$  matrices.
- Q14.** `model` on CPU, `x=randn(32,784).to("cuda")`, then `model(x)`.
- Q15.** Training loop missing one line — gradients accumulate across batches.
- Q16.** `probs=Softmax(logits)`, then `loss=CrossEntropyLoss(probs,y)`.
- Q17.** Eval loop has no `torch.no_grad()` wrapper.

### Section D — Math Questions

- Q18.**  $y = 3x^3 + 2x + 1$ : find  $dy/dx$  at  $x=4$ . What does `x.grad` print?
- Q19.** `Linear(512,256)`: how many trainable parameters?
- Q20.** `Conv2d(3,64,3)`: how many parameters?
- Q21.** PIL image ( $H=256, W=256, C=3$ ): what converts to PyTorch format?
- Q22.** Output  $(32, 10)$ : write 2 lines for predictions and accuracy vs  $y$  shape  $(32, )$ .

## SOLUTIONS

### Section A

- A1.** Static graphs compile before execution (fixed). Dynamic build during execution — normal Python control flow, `print()`, `pdb` all work directly.
- A2.** `dtype` (data type), `shape` (dimensions), `device` (`cpu/cuda/mps`).
- A3.** Gradients *accumulate* (add to `.grad`). Without clearing, batches 2+ have wrong accumulated gradients baked into updates.
- A4.** `CrossEntropyLoss` applies log-softmax internally. Passing softmax outputs doubles it — wrong loss, no error thrown. Silent bug.
- A5.** `train()`: Dropout active, BN uses batch stats. `eval()`: Dropout off, BN uses running stats accumulated during training.

$$\text{A6. } \frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial f} \cdot \frac{\partial f}{\partial g} \cdot \frac{\partial g}{\partial h} \cdot \frac{\partial h}{\partial \mathbf{x}}$$

**A7.** Tells PyTorch to track all ops on this tensor in a DAG. Graph built *dynamically* during forward pass. `.backward()` traverses in reverse.

### Section B

- B6.** `view(2)?` `view(4,-1)`:  $(4, 2352)$  since  $3 \times 28^2 = 2352$ ; `permute(0,2,3,1)`:  $(4, 28, 28, 3)$ ; `[:,0,:,:]`:  $(4, 28, 28)$
- B9.**  $(32, 128)$  — inner dims match; outer dims form result.

**B10.**  $(8, 16) \rightarrow (1, 8, 16) \rightarrow (1, 8, 1, 16)$

**B11.**  $(3, 5)$  — size-1 middle dim removed.

**B12.** `t.unsqueeze(1)` or `t.reshape(100,1)`

### Section C

- C13.**  $x * y = \text{element-wise}$ . For matmul: `x @ y` or `torch.matmul(x,y)`.

**C14.** Model is on CPU, data on CUDA. Fix: `model.to("cuda")` before `model(x)`.

**C15.** Missing `optimizer.zero_grad()` before `loss.backward()`.

**C16.** Softmax before CrossEntropyLoss — double applies. Pass raw logits.

**C17.** Wrap with `with torch.no_grad():` — saves memory, faster.

### Section D

**D18.**  $dy/dx = 9x^2 + 2$ . At  $x=4$ : **146**. `x.grad = tensor(146.)`

**D19.**  $(512 \times 256) + 256 = \mathbf{131,328}$

**D20.**  $64 \times (3 \times 9 + 1) = 64 \times 28 = \mathbf{1,792}$

**D21.** `tensor.permute(2,0,1): (H, W, C) \rightarrow (C, H, W)`

**D22.** `preds=output.argmax(dim=1)`  
`acc=(preds==y).float().mean().item()`

## TIPS & TRICKS — LOCK IN

- Shape debug first.** Print `t.shape` at every step. 90% of errors are shape mismatches.
- Always `loss.item()` to log** — avoids accumulating GPU tensors.
- Normalise inputs to  $[-1, 1]$ :** use  $(x - \mu)/\sigma$ .
- Default:** Adam, `lr=1e-3`. Too high: loss explodes. Too low: stuck.
- Set seed:** `torch.manual_seed(42)` for reproducibility.
- No softmax before CrossEntropyLoss.** Silent error, not a crash.
- model.parameters()** for optimizer; `state_dict()` for saving.
- no\_grad()** always during inference — faster, less memory.
- Classification predictions:** `output.argmax(dim=1)`.
- Start small.** Debug on 10 samples, then scale up.

## QUICK REFERENCE

Concept	Syntax
Matrix multiply	<code>A @ B</code>
Enable gradient	<code>requires_grad=True</code>
Backward pass	<code>loss.backward()</code>
Zero gradients	<code>optimizer.zero_grad()</code>
Weight update	<code>optimizer.step()</code>
Stop gradients	<code>torch.no_grad()</code>
Move to GPU	<code>tensor.to(device)</code>
Predictions	<code>out.argmax(dim=1)</code>
Save model	<code>torch.save(m.state_dict(),...)</code>
Train mode	<code>model.train()</code>
Eval mode	<code>model.eval()</code>

### The Sacred 5-Step Training Loop

- |  |  |
|--|--|
| 1. <code>pred=model(X)</code><br>2. <code>loss=crit(pred,y)</code><br>3. <code>optimizer.zero_grad()</code><br>4. <code>loss.backward()</code><br>5. <code>optimizer.step()</code> | forward<br>loss<br>clear grads<br>backprop<br>update weights |
|--|--|